

Homework Assignment 1

Due Wednesday, Sep. 19, at 9:00 am

OVERVIEW

The purpose of this assignment is to introduce you to Matlab as a tool for manipulating images. For this, you will build your own version of a very basic image processing pipeline (without denoising and color transform steps), to convert a RAW image into an image that can be displayed on a computer monitor or printed on paper. There is a “Hints and Information” section at the end of this document that is likely to help.

1. IMPLEMENT A BASIC IMAGE PROCESSING PIPELINE (90 POINTS)

Throughout this problem, you will use the file `banana slug.tiff` included in the `./data` directory of the homework ZIP archive. This is a RAW image that was captured with a Canon EOS T3 Rebel camera. We did some very slight pre-processing to the original RAW file, in order to convert it to `.tiff` format. At the end of this problem, the image should look something like what is shown in Figure 1. The exact result can vary greatly, depending on the choices you make in your implementation of the image processing pipeline.



Figure 1: One possible rendering of the RAW image provided with this assignment.

INITIALS (5 PTS)

Load the image into Matlab. Originally, it will be in the form of a 2D-array of unsigned integers. Check and report how many bits per integer the image has, and what its width and height is. Then, convert the image into a double-precision array. (See help for functions `imread`, `size`, `class` and `double`.)

LINEARIZATION (5 POINTS)

The 2D-array is not yet a linear image. As we discussed in class, it is possible that it has an offset due to dark noise, and saturated pixels due to overexposure. Additionally, even though the original data-type of the image was 16 bits, only 14 of those have meaningful information, meaning that the maximum possible value for pixels is 16383 (that's $2^{14} - 1$). For the provided image file, you can assume the following: All pixels with a value lower than 2047 correspond to pixels that would be black, were it not for dark noise. All pixels with a value above 15000 are overexposed pixels. (The values 2047 for the black level and 15000 for saturation are taken from the manufacturer).

Convert the image into a linear array within the range $[0, 1]$. Do this by applying a linear transform (shift and scale) to the image, so that the value 2047 is mapped to 0, and the value 15000 is mapped to 1. Then, clip negative values to 0, and values greater than 1 to 1. (See help for functions `min` and `max`.)

IDENTIFYING THE CORRECT BAYER PATTERN (20 POINTS)

As we discussed in class, most cameras use the Bayer pattern in order to capture color. The same is true for the camera used to capture our RAW image. We do not know, however, the exact shift of the Bayer pattern relative to our image: If you look at the top-left 2×2 square of the image file, it can correspond to any of four possible red-green-blue patterns, as shown in Figure 2.




Figure 2: From left to right: 'grbg', 'rggb', 'bggr', 'gbrg'.

Think of a way for identifying which version of the Bayer patterns applies to our image file, and report which version you identified. (See Hints and Information)

WHITE BALANCING (20 POINTS)

After identifying the correct Bayer pattern, we want to do white balancing. Implement both the white world and gray world automatic white balancing algorithms, as discussed in class. At the end of this assignment, check what the image looks like under both white balancing algorithms, and decide which one you like best. (See help for function `mean`)

DEMOSAICING (25 POINTS)

After white balancing, you want to demosaic the image. Use bilinear interpolation for demosaicing, as discussed in class. Do not implement bilinear interpolation manually. Instead, read the document to learn how to use Matlab function `interp2` function for this purpose.

BRIGHTNESS ADJUSTMENT AND GAMMA CORRECTION (20 POINTS)

You now have a 16-bit, full-resolution, linear RGB image. Because of the scaling you did at the start of the assignment, the image pixel values may not be in a range appropriate for display. Additionally, as we discussed in class, the image is not yet gamma-corrected. As a result, when displaying the image, it will appear very dark.

Brighten the image by linearly scaling it by some number. Select the scale as a percentage of the pre-brightening maximum grayscale value. (see help for `rgb2gray` for converting the image to grayscale). The correct percentage is highly subjective, so you should experiment with many different percentages and report and use what percentage looks best to you.

Before having an image that can be properly displayed, the last step you need to do is tone reproduction (gamma correction). For this, implement the following non-linear transform, then apply it to the image:

$$C_{\text{non-linear}} = \begin{cases} 12.92 \cdot C_{\text{linear}}, & C_{\text{linear}} \leq 0.0031308 \\ (1 + 0.055) \cdot C_{\text{linear}}^{\frac{1}{2.4}} - 0.055, & C_{\text{linear}} \geq 0.0031308 \end{cases}$$

Where $C = \{R, G, B\}$ is each of the red, green, and blue channels. This function may look completely arbitrary but it comes from the sRGB standard. It is a good default choice if the camera's true gamma correction curve is not known.

COMPRESSION (5 PTS) Finally, it is time to store the image, either with or without compression. Use the `imwrite` command to store the image in .PNG format (no compression), and also in .JPEG format with quality setting 95. This setting determines the amount of compression. Can you tell the difference between the two files? The compression ratio is the ratio between the size of the uncompressed file (in bytes) and the size of the compressed file (in bytes). What is the compression ratio?

By changing the JPEG quality settings, determine the lowest setting for which the compressed image is indistinguishable from the original. What is the compression ratio?

DELIVERABLES

For this project, and all other projects, you must do a project report in HTML. In the report you will describe your algorithm and any decisions you made to write your algorithm a particular way. Use both words and images to show us what you've done (describe in detail your algorithm parameterization for each of your results). You should submit your code and the link for HTML by the deadline.

HINTS AND INFORMATION

- To get help on a particular function in Matlab, type `help <function>`. To get a list of functions related to a particular keyword, use the `lookfor` function. We will be making extensive use of the Image Processing Toolbox, and a list of the functions in that toolbox is generated by typing `help images`. Print your results (to a printer or to a file) using the `print` command, and when making hardcopies please save space by using `subplot` whenever possible. As an example, the following Matlab script loads three images, displays them in a figure and prints the figure to a PNG file.

```
% read three images from current directory
im1 = imread('image1.tiff');
im2 = imread('image2.tiff');
im3 = imread('image3.tiff');

% display images in a figure, side-by-side
figure; % create a new figure
imshow(im1); % display an image
title('Image 1'); % add a title

% print the displayed figure a PNG file. You can also print
from the figure menubar.
print -dpng output.png
```

- The colon operator `:` allows to form arrays out of subsets of other arrays. In the following example, given an original image `im`, it creates three other images, each with only one-fourth the pixels of the originals. The pixels of each of the corresponding sub-images are shown in Figure 3. You can also use the function `cat` to combine these three images into a single 3-channel RGB image.

```
% create three sub-images of im as shown in figure below
im1 = im(1:2:end, 1:2:end)
im2 = im(1:2:end, 2:2:end);
```

```

im3 = im(2:2:end, 1:2:end);
% combine the above images into an RGB image, such
that im1 is the red,
% im2 is the green, and im3 is the blue channel
im_rgb = cat(3, im1, im2, im3);

```

- You will find it very helpful to display intermediate results while you are implementing the image processing pipeline. However, before you apply the brightening and gamma correction, you will find that displayed images will look completely black. To be able to see something more meaningful, you can use the following command to display an intermediate image `im_intermediate`.

```
figure; imshow(min(1, im_intermediate * 5));
```

For example, Figure 3 shows what you should see using this command after correctly performing the linearization step.

Im1	Im2	Im1	Im2	Im1	Im2
Im3		Im3		Im3	
Im1	Im2	Im1	Im2	Im1	Im2
Im3		Im3		Im3	
Im1	Im2	Im1	Im2	Im1	Im2
Im3		Im3		Im3	



Figure 3: Left: The linear RAW image (brightness increased by 5). Right: Crop for showing the Bayer pattern.

POLICY

Do not search the code on the web. Once detected, you will be penalized by – 200% on your project credits.